

Quantitative Macroeconomics w/ AI and ML

Lec. 3: Programming Basics for Economists

Zhigang Feng

Dec., 2025

Overview

Today's Goal

By the end of this lecture, you should be able to:

- Understand the **economist's tech stack**: OS, Python, IDEs, libraries, and AI assistants (Cursor).
- Read and write basic **Python** code:
 - data types, lists, dictionaries, loops, functions, classes.
- Follow a typical **data science workflow**:
 - download, clean, save, and visualize data (e.g., FRED GDP series).
- Get a **first flavor** of **PyTorch** and the ML workflow:
 - tiny neural net: predict next-quarter GDP growth from last 4 quarters.
 - ML theory is next lecture; today is about *how the code fits together*.

The Economist's Tech Stack

How the Pieces Fit Together

- 1. **Operating System (OS)**
 - Windows, macOS, Linux.
 - Manages hardware (CPU, RAM, file system).
 - Python and all tools run *on top* of this.
- 2. **Python**
 - High-level programming language for data analysis, simulations, and ML.
 - You install it (e.g., via **Anaconda** or [python.org](https://www.python.org)).
- 3. **IDE / Notebook**
 - Where you edit and run code:
 - **Jupyter**, **VS Code/Cursor**, or **PyCharm**.

Jupyter, VS Code/Cursor, and PyCharm

Jupyter (Notebooks)

- Run code **cell-by-cell**.
- Great for:
 - exploration,
 - data cleaning,
 - plotting.
- Immediate feedback: results appear under each cell.
- Comes with Anaconda as **JupyterLab**.
- Recommended for:
 - homework,
 - small experiments,
 - demonstrations in this course.

VS Code/Cursor and PyCharm

- **VS Code/Cursor:**
 - general-purpose editor/IDE,
 - becomes a full Python IDE with the Python extension,
 - excellent for both **small** and **large** projects.
- **PyCharm:**
 - dedicated **Python IDE**,
 - more Python-specific tools out of the box (refactoring, inspections),
 - great for bigger research codebases.
- Both can handle complex projects. Choice is about personal taste and features.

Python, PyTorch, and AI: How They Relate

- **Python:**

- General-purpose language.
- You use it for **data cleaning, regression, simulations, plotting.**

- **PyTorch:**

- A **Python library** specialized for:
 - tensors (multi-dimensional arrays),
 - automatic differentiation (autograd),
 - neural networks (deep learning).

- **AI / ML Models:**

- Built in **PyTorch using Python.**
- Examples:
 - time series forecasters,
 - macro-finance deep learning models,
 - text models for central bank communications.

Takeaway

Python is the language. PyTorch is the toolbox. AI / ML models are what you *build* with that toolbox.

AI Assistants in Your Stack (Cursor, Copilot, . . .)

- **AI Coding Tools:**

- Cursor, GitHub Copilot, ChatGPT-style tools.

- Think of them as **supercharged TAs**:

- complete boilerplate code,
- explain error messages,
- suggest better structure,
- translate math into code.

- Use them especially when:

- installation / environment is broken,
- you see a mysterious error,
- you are implementing a new algorithm from a paper.

Using Cursor Effectively for This Class

- **Cursor** = VS Code-like editor with AI deeply integrated.
- Helpful modes:
 - **Inline suggestions**: completes loops, function definitions.
 - **Chat panel**: ask questions with your code context.
 - **Explain / Fix**: highlight code → ask “Explain this” or “Fix this error”.

Example Prompts in Cursor

Debugging:

I'm trying to run this PyTorch training loop, but I get a shape mismatch error.

Here is the full code and error. Explain what's wrong and show a fixed version.

Code comprehension:

Explain this class in simple terms as if I were an economics student.

What are the inputs, outputs, and what does forward() do?

Environment help:

My Python environment has both conda and pip. How should I install pandas and torch correctly on macOS?

Installing Python and Tools

Anaconda, Python, and Editors: Who Does What?

- **Anaconda:**

- Distribution that installs:
 - Python interpreter,
 - many scientific packages (NumPy, pandas, etc.),
 - environment manager (`conda`),
 - **JupyterLab** as a notebook interface.
- It does *not* replace an editor: you still need a place to write code.

- **Editors/IDEs:**

- **JupyterLab** (browser) — comes with Anaconda.
- **VS Code/Cursor** — external editor; can use Anaconda's Python.
- **PyCharm** — external Python IDE; can also use Anaconda envs.

Mental model

Anaconda = **Python + packages + env manager**.

VS Code/Cursor/PyCharm = **where you write, run, and debug your code**.

Installing Python (Practical Version)

- **Easiest path (recommended): Anaconda**

- Download from <https://www.anaconda.com/download>
- Includes Python + many scientific packages (NumPy, pandas, etc.).

- **Alternative: python.org**

- Download from <https://www.python.org/downloads/>
- On Windows: check “Add Python to PATH”.

- **Verify installation** in terminal / command prompt:

```
1 python --version      # or python3 --version
```

- If this fails: **ask AI (Cursor / ChatGPT) for help** using a clear prompt and include the full error.

Managing Packages

- Use pip (or conda if using Anaconda) to install libraries:

```
1 # Using pip
2 pip install pandas pandas-datareader matplotlib seaborn torch
3
4 # Using conda (Anaconda)
5 conda install pandas matplotlib seaborn
6 conda install -c pytorch pytorch
```

- You install each package **once per environment**, not every time you run the code.
- If installation errors appear, copy the error text and ask:
 - a classmate,
 - AI (with full error),
 - or the instructor.

Python Basics: Data Types & Structures

Atomic Data Types

Smallest units of data in Python:

- **Integers (int)**: whole numbers.

```
1 year = 2025
```

- **Floats (float)**: real numbers (decimals).

```
1 inflation_rate = 0.035
```

- **Strings (str)**: text.

```
1 country = "United States"
```

- **Booleans (bool)**: logical truth values.

```
1 is_recession = False
```

Working with Text (Strings)

Economists often clean country names, tickers, etc.

```
1 country_raw = "    united states    "
2 gdp = 23.5
3
4 # f-strings: modern way to format text
5 print(f"The GDP of {country_raw} is {gdp} trillion.")
6 # Output: "The GDP of    united states    is 23.5 trillion."
7
8 # Cleaning methods
9 country_clean = country_raw.strip().title()
10 print(country_clean)
11 # Output: "United States"
```

Lists: Ordered Collections

Lists behave like columns of numbers in Excel.

```
1 # Quarterly GDP growth (%)
2 gdp_growth = [1.2, 2.5, 3.1, -0.5]
3
4 # Indexing (starts at 0)
5 first = gdp_growth[0]      # 1.2
6
7 # Slicing
8 subset = gdp_growth[1:3]    # [2.5, 3.1]
9
10 # Modifying
11 gdp_growth[3] = 0.1        # replace -0.5 with 0.1
```

Warning: Reference vs Copy

Important: variable names are just *labels*, not boxes.

```
1 a = [1, 2, 3]
2 b = a          # b points to the SAME list
3 b[0] = 99
4
5 print(a)      # [99, 2, 3]
```

To make an independent copy:

```
1 c = a.copy()  # now c is separate
```

This idea (sharing vs copying) also appears with **NumPy arrays** and **PyTorch tensors**.

Dictionaries: Key-Value Data

Dictionaries map **keys** to **values**, like labeled columns.

```
1 country_data = {  
2     "name": "Japan",  
3     "gdp": 4.2,           # trillion USD  
4     "population": 125,    # million  
5     "currency": "Yen"  
6 }  
7  
8 print(country_data["gdp"])        # 4.2  
9 country_data["income_per_cap"] = (  
10    country_data["gdp"] * 1e12 / (125e6)  
11 )
```

Think of this as a tiny JSON object. Pandas DataFrames are like **tables of dictionaries**.

Tuples and Sets

- **Tuples ()**: like lists, but **immutable**.

```
1 coords = (41.25, -95.93)    # (lat, lon)
```

Useful for fixed data (e.g., parameters, coordinates, (year, quarter)).

- **Sets {}**: unordered collection of **unique** items.

```
1 industry_codes = ["A", "B", "A", "C"]
2 unique_codes = set(industry_codes)    # {"A", "B", "C"}
```

Great for removing duplicates or checking membership.

Python Basics: Logic, Functions, Modules

Operators

- **Arithmetic:** +, -, *, /
- Power: ** (e.g., $2^{**3} = 8$)
- Modulo: % (remainder, e.g., $17\%4 = 1$)
- **Comparison:** ==, !=, >, <, >=, <=
- **Logical:** and, or, not

Conditionals and Loops

If / Elif / Else

```
1 x = -1
2
3 if x > 0:
4     print("Positive")
5 elif x < 0:
6     print("Negative")
7 else:
8     print("Zero")
```

For Loop

```
1 gdp_growth = [1.2, 2.5, -0.5]
2
3 for g in gdp_growth:
4     if g < 0:
5         print("Recession quarter!")
```

While Loop

```
1 count = 0
2 while count < 3:
3     print(count)
4     count += 1
```

List Comprehensions

Compact way to create lists. Like set-builder notation.

- Math: $S = \{x^2 \mid x \in \mathbb{N}, x > 0\}$
- Python:

```
1 data = [1, 2, 3, 4, 5]
2
3 # Old way
4 squares = []
5 for x in data:
6     squares.append(x**2)
7
8 # Pythonic way
9 squares = [x**2 for x in data]
10
11 # Filter + transform
12 gdp_growth = [1.2, 2.5, -0.5, 0.3]
13 positive_growth = [g for g in gdp_growth if g > 0]
```

Error Handling: try / except

Economic data is messy. Use try/except to avoid crashes.

```
1 values = [100, 200, 0, 300]
2
3 for v in values:
4     try:
5         ratio = 10 / v
6         print(f"Ratio: {ratio}")
7     except ZeroDivisionError:
8         print("Error: Division by zero, skipping...")
```

You can also catch ValueError when parsing strings from CSV files.

Functions: Reusable Blocks of Logic

```
1 def cobb_douglas(k, l, alpha=0.3):
2     """Output Y = K^alpha * L^(1-alpha)"""
3     y = k**alpha * l**(1 - alpha)
4     return y
5
6 # Usage
7 output = cobb_douglas(k=100, l=50)
8 print(output)
```

- **Inputs** are arguments (k, l, alpha).
- **Outputs** come from return.
- Docstring """"...""" explains the function.

Modules and Imports

- **Module:** a .py file containing Python code.
- **Package:** collection of modules (e.g., pandas, torch).

```
1 import math  
2 print(math.sqrt(16))      # 4.0  
3  
4 import numpy as np  
5 import pandas as pd
```

Custom module example:

```
1 # my_module.py  
2 def square(x):  
3     return x * x  
4  
5 # main.py  
6 import my_module  
7 print(my_module.square(5))  # 25
```

Organization: Classes and Objects

Why Economists Need OOP

- **State management:**

- An agent has wealth, age, employment status, expectations.
- A class bundles these attributes together.

- **Behavior:**

- Agent can consume, work, update beliefs, etc.
- Methods define these behaviors.

- **PyTorch models are classes:**

- A neural network is a class with:
 - attributes: layers (parameters),
 - methods: `forward` (how to compute outputs).

Anatomy of a Class

```
1 class Agent:
2     def __init__(self, wealth, age):
3         # self = "this instance"
4         self.wealth = wealth    # attribute
5         self.age     = age       # attribute
6
7     def consume(self, amount):
8         # method (behavior)
9         self.wealth -= amount
10
11    def birthday(self):
12        self.age += 1
13
14 # Instantiate (create objects)
15 a1 = Agent(wealth=1000, age=30)
16 a2 = Agent(wealth=500,   age=25)
17
18 # Call methods
19 a1.consume(100)
20 a1.birthday()
```

Without Classes vs With Classes

Without classes (state scattered):

```
1 # Separate variables
2 wealth_1 = 1000
3 age_1    = 30
4
5 wealth_2 = 500
6 age_2    = 25
7
8 def consume(wealth, amount):
9     return wealth - amount
10
11 wealth_1 = consume(wealth_1, 100)
12 age_1    = age_1 + 1
```

With classes (state + behavior bundled):

```
1 class Agent:
2     def __init__(self, wealth, age):
3         self.wealth = wealth
4         self.age   = age
```

Inheritance and Why PyTorch Uses It

```
1 # Parent class
2 class GenericModel:
3     def predict(self, x):
4         print("Predicting... ")
5
6 # Child class inherits from parent
7 class MyModel(GenericModel):
8     pass
9
10 m = MyModel()
11 m.predict() # "Predicting... "
```

- In PyTorch, your model will inherit from `nn.Module`:
 - registers parameters automatically,
 - integrates with optimizers and device management,
 - provides `model.parameters()`, `model.to(device)`, etc.

Data Science for Economists

The Data Workflow

Typical pipeline in an empirical macro project:

1. **Acquire** data (FRED, World Bank, CSVs).
2. **Clean**:
 - handle missing values, correct types, adjust dates.
3. **Persist**:
 - save clean data (CSV / pickle) so you do not redo steps 1–2.
4. **Analyze / Visualize**:
 - regressions, summary statistics, plots.
5. **Export**:
 - figures and tables for your LaTeX paper.

Acquire: FRED GDP Data

```
1 import pandas_datareader.data as web
2 import datetime
3
4 start = datetime.datetime(1990, 1, 1)
5 end   = datetime.datetime(2025, 1, 1)
6
7 # 'GDP' is the FRED series ID for real GDP (quarterly)
8 gdp_data = web.DataReader('GDP', 'fred', start, end)
9
10 print(gdp_data.head())
```

- `pandas_datareader` sends a web request to the FRED API.
- FRED identifies each series by its **series ID** (here: '`GDP`').
- The response (dates + values) is converted into a **pandas DataFrame**.
- Requires an internet connection. For heavy use, you can obtain an API key from FRED, but many simple calls work without.

Clean: Compute Growth and Filter

```
1 import pandas as pd
2 import numpy as np
3
4 # Compute quarterly growth rate
5 gdp_data['growth'] = gdp_data['GDP'].pct_change()
6
7 # Drop missing values from the first difference
8 clean = gdp_data.dropna()
9
10 # Example: post-2000
11 post_2000 = clean[clean.index >= '2000-01-01']
12
13 print(post_2000[['GDP', 'growth']].head())
```

Example: HP Filter (Trend and Cycle)

```
1 import numpy as np
2 from statsmodels.tsa.filters.hp_filter import hpfilter
3
4 # Log GDP
5 gdp_log = np.log(df['GDP'])
6
7 # HP filter with lambda=1600 for quarterly data
8 cycle, trend = hpfilter(gdp_log, lamb=1600)
9
10 df['gdp_trend'] = trend
11 df['gdp_cycle'] = cycle
```

```
1 import matplotlib.pyplot as plt
2
3 plt.figure(figsize=(10, 5))
4 plt.plot(df.index, gdp_log, label="log GDP", alpha=0.5)
5 plt.plot(df.index, df['gdp_trend'], label="HP trend", linewidth=2)
6 plt.title("HP Filter: log GDP and Trend")
7 plt.xlabel("Year")
8 plt.legend()
```

Persist: Saving Clean Data

```
1 # Save to CSV for sharing or inspection
2 post_2000.to_csv("gdp_post2000.csv")
3
4 # Save to pickle (Python-specific, preserves types)
5 post_2000.to_pickle("gdp_post2000.pkl")
6
7 # Later: load from CSV
8 df = pd.read_csv("gdp_post2000.csv",
9                  index_col='DATE', parse_dates=True)
```

Visualize: GDP Growth

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3
4 sns.set_theme(style="whitegrid")
5
6 x = df.index
7 y = df['growth'] * 100 # to percent
8
9 plt.figure(figsize=(10, 5))
10 plt.plot(x, y, label="Quarterly growth (%)")
11 plt.axhline(0, linestyle='--')
12 plt.title("US Real GDP Growth")
13 plt.xlabel("Year")
14 plt.ylabel("Growth (%)")
15 plt.legend()
16 plt.tight_layout()
```

Export: High-Quality Figures

```
1 # Save for your paper
2 plt.savefig("gdp_growth.png", dpi=300,
3             bbox_inches='tight')
4
5 # Vector graphics (good for LaTeX)
6 plt.savefig("gdp_growth.pdf", bbox_inches='tight')
7
8 plt.show()
```

PyTorch Basics

What is PyTorch?

- Open-source deep learning library in Python.
- Key components:
 - **Tensors**: multi-dimensional arrays (like NumPy arrays).
 - **Autograd**: automatic differentiation.
 - **nn.Module**: building neural networks.
 - **optim**: optimizers (SGD, Adam, etc.).
- Widely used for:
 - time series forecasting,
 - macro / finance ML,
 - text and image models.

PyTorch: Applications and Advantages

- **Applications:**

- Computer vision: image classification, object detection.
- NLP: text classification, language models, machine translation.
- Reinforcement learning: training agents to interact with environments.
- Economics: forecasting, structural models with neural nets, text analysis of policy statements.

- **Advantages:**

- **Dynamic computation graph:**

- Graph is built at runtime, easy to debug and modify.

- **GPU acceleration:**

- seamless move between CPU and GPU for speed.

- **Active community:**

- many tutorials, extensions, and research codebases.

Tensors: Like NumPy, but GPU-Ready

```
1 import torch
2
3 # 1D tensor (vector)
4 t1 = torch.tensor([1, 2, 3, 4, 5])
5 print(t1)
6
7 # 2D tensor (matrix)
8 t2 = torch.tensor([[1, 2, 3],
9                   [4, 5, 6]])
10 print(t2)
11
12 # Specify dtype
13 t_float = torch.tensor([1.0, 2.0, 3.0],
14                      dtype=torch.float32)
15
16 # Zeros
17 t_zeros = torch.zeros((3, 4))
```

Tensors and NumPy arrays can be converted back and forth.

Tensors (Visual)


1 D TENSOR /
VECTOR

5
7
4 5
1 2
- 6
3
2 2
1
6
3
- 9


2 D TENSOR /
MATRIX

- 9	4	2	5	7
3	0	1 2	8	6 1
1	2 3	- 6	4 5	2
2 2	3	- 1	7 2	6


3 D TENSOR /
CUBE



- 9	4	2	5	7
3	0	1 2	8	6 1
1	2 3	- 6	4 5	2
2 2	3	- 1	7 2	6



4 D TENSOR
VECTOR OF CUBES



5 D TENSOR
MATRIX OF CUBES

Tensors: Summary

- Multi-dimensional arrays with a uniform data type:
 - e.g., all float32 or all int64.
- Similar to NumPy arrays:
 - support element-wise operations, reshaping, indexing, slicing.
 - can be moved between CPU and GPU memory.
- Designed for efficient numerical computation:
 - can handle large matrices and high-dimensional data.

Basic Tensor Operations

```
1 x = torch.tensor([1, 2, 3])
2 y = torch.tensor([4, 5, 6])
3
4 sum_result      = x + y    # [5, 7, 9]
5 product_result = x * y    # [4, 10, 18]
6
7 # Matrix-style operations
8 A = torch.tensor([[1.0, 2.0],
9                  [3.0, 4.0]])
10 B = torch.tensor([[5.0, 6.0],
11                   [7.0, 8.0]])
12
13 C = A @ B  # matrix multiplication
```

Reshaping: view() vs reshape()

```
1 x = torch.arange(1, 10)      # [1, 2, ..., 9]
2
3 # Want 3x3
4 y = x.view(3, 3)            # requires x to be contiguous
5 z = x.reshape(3, 3)          # more flexible
6
7 print(y)
8 # tensor([[1, 2, 3],
9 #         [4, 5, 6],
10 #        [7, 8, 9]])
```

- `view`: returns a new tensor that shares memory, but only works if memory layout is contiguous (otherwise you may need `x = x.contiguous().view(...)`).
- `reshape`: tries to return a view; if that fails, it creates a copy.
- In most beginner code, `reshape` is a safe default.

Concatenation and Transpose

```
1 x = torch.tensor([[1, 2],  
2                 [3, 4]])  
3 y = torch.tensor([[5, 6],  
4                 [7, 8]])  
5  
6 # Concatenate vertically (rows)  
7 cat_rows = torch.cat((x, y), dim=0)  
8  
9 # Concatenate horizontally (cols)  
10 cat_cols = torch.cat((x, y), dim=1)  
11  
12 # Transpose  
13 xt = x.T    # or torch.transpose(x, 0, 1)
```

Autograd: Conceptual View


- PyTorch has an **automatic differentiation** engine: autograd.
- When you operate on tensors with `requires_grad=True`:
 - PyTorch builds a **computation graph** where:
 - nodes = tensors,
 - edges = operations.
- When you call `backward()` on a scalar output (typically a loss):
 - PyTorch traverses the graph backward,
 - applies the chain rule,
 - computes gradients for all tensors that require them.
- These gradients are stored in the `.grad` attribute of tensors (e.g., model parameters).

Autograd: Simple Example

```
1 import torch
2
3 x = torch.tensor(2.0, requires_grad=True)
4 y = x * x                      # y = x^2
5
6 y.backward()                     # dy/dx at x=2
7 print(x.grad)                  # tensor(4.)
```

- Here, y is the **loss** (a scalar).
- `backward()` computes $\frac{dy}{dx}$.
- In neural nets, parameters (weights, biases) play the role of x .

Neural Networks and Gradient Descent (Visual)



Autograd: Steps for Optimization

1. **Enable gradient tracking:** create tensors with `requires_grad=True` (PyTorch does this for parameters in `nn.Module`).
2. **Forward pass:** compute model output and loss.
3. **Backward pass:** call `loss.backward()` to compute gradients.
4. **Update parameters:** use optimizer (e.g., `optimizer.step()`).
5. **Zero gradients:** call `optimizer.zero_grad()` before the next iteration.


Optimizers and Training Loop Template

```
1 import torch.nn as nn
2 import torch.optim as optim
3
4 model = nn.Linear(4, 1)          # simple linear model
5 optimizer = optim.SGD(model.parameters(), lr=0.01)
6 loss_fn = nn.MSELoss()
7
8 num_epochs = 10
9 for epoch in range(num_epochs):
10     # x_batch, y_batch = ...
11     optimizer.zero_grad()        # 1. clear old gradients
12     preds = model(x_batch)      # 2. forward
13     loss = loss_fn(preds, y_batch) # 3. loss
14     loss.backward()              # 4. backprop
15     optimizer.step()            # 5. update params
16
17     print(f"Epoch {epoch+1}, loss = {loss.item():.4f}")
```

Steps of a DNN Training Workflow (Reference)

1. Import libraries: `torch`, `torch.nn`, `torch.optim`, `datasets`.
2. Prepare data:
 - load or generate `x` and `y`,
 - convert to PyTorch tensors.
3. Build a dataset class (optional) and wrap with `DataLoader`.
4. Define the neural network as a class inheriting from `nn.Module`.
5. Choose loss function (e.g., MSE) and optimizer (e.g., Adam).
6. Train:
 - loop over epochs and batches,
 - forward → loss → backward → step.
7. Evaluate on validation/test data and visualize results.

Mini-batches and DataLoader (Visual)



DataLoader handles batching, shuffling, and iteration over the dataset.

DataLoader: Mini-batches (Concept)

- `torch.utils.data.DataLoader`:
 - splits data into mini-batches,
 - can shuffle each epoch,
 - can load data in parallel (via `num_workers`).
- Properties:
 - Each epoch: all data points used once (unless `drop_last=True`).
 - `batch_size`: trade-off between speed and stability.
 - `shuffle=True`: helpful for training neural nets.

DataLoader Example

```
1 from torch.utils.data import Dataset, DataLoader
2 import torch
3
4 class CustomDataset(Dataset):
5     def __init__(self, data, labels):
6         self.data = data
7         self.labels = labels
8
9     def __len__(self):
10        return len(self.data)
11
12    def __getitem__(self, idx):
13        return self.data[idx], self.labels[idx]
14
15 dataset = CustomDataset(torch.randn(1000, 10),
16                         torch.randint(0, 2, (1000,)))
17 loader = DataLoader(dataset, batch_size=32,
18                     shuffle=True)
19
20 for batch_x, batch_y in loader:
21     print(f'batch_x shape: {batch_x.shape}, batch_y shape: {batch_y.shape}')
```

Monitoring Progress with tqdm

```
1 from tqdm import tqdm
2
3 for epoch in range(num_epochs):
4     progress = tqdm(loader,
5                       desc=f"Epoch {epoch+1}/{num_epochs}")
6     for batch_x, batch_y in progress:
7         optimizer.zero_grad()
8         preds = model(batch_x)
9         loss = loss_fn(preds, batch_y)
10        loss.backward()
11        optimizer.step()
12        progress.set_postfix({"loss": loss.item()})
```

tqdm provides a progress bar and live loss updates.

GPU Acceleration (Optional for This Class)

```
1 device = torch.device("cuda" if torch.cuda.is_available()
2                         else "cpu")
3
4 model = model.to(device)
5
6 for batch_x, batch_y in loader:
7     batch_x = batch_x.to(device)
8     batch_y = batch_y.to(device)
9     # training steps...
```

- For this course, **CPU is fine** for small examples.
- GPU becomes important for:
 - large datasets,
 - big neural networks.

Reference: A Deeper MLP with Dropout

```
1 import torch.nn as nn
2
3 class EqumNN(nn.Module):
4     def __init__(self, n1, n2, dropout_prob=0.0):
5         super().__init__()
6         self.fc1 = nn.Linear(n1, 64)
7         self.dropout1 = nn.Dropout(p=dropout_prob)
8         self.fc2 = nn.Linear(64, 32)
9         self.dropout2 = nn.Dropout(p=dropout_prob)
10        self.fc3 = nn.Linear(32, n2)
11
12    def forward(self, x):
13        x = nn.functional.relu(self.fc1(x))
14        x = self.dropout1(x)
15        x = nn.functional.relu(self.fc2(x))
16        x = self.dropout2(x)
17        x = self.fc3(x)
18
19        return x
```

EqumNN: Mathematical Representation

Let $x \in \mathbb{R}^{n_1}$ be the input and $y_3 \in \mathbb{R}^{n_2}$ the output.

- First layer:

$$y_1 = \text{ReLU}(W_1x + b_1), \quad W_1 \in \mathbb{R}^{64 \times n_1}, \quad b_1 \in \mathbb{R}^{64}.$$

- Second layer:

$$y_2 = \text{ReLU}(W_2y_1 + b_2), \quad W_2 \in \mathbb{R}^{32 \times 64}, \quad b_2 \in \mathbb{R}^{32}.$$

- Third layer:

$$y_3 = W_3y_2 + b_3, \quad W_3 \in \mathbb{R}^{n_2 \times 32}, \quad b_3 \in \mathbb{R}^{n_2}.$$

- Dropout is used during training between layers as a regularization technique (randomly zeroes some components of y_1 and y_2).

In summary:

$$y_3 = W_3 \text{ReLU}(W_2 \text{ReLU}(W_1x + b_1) + b_2) + b_3.$$

Mini Example: Predict GDP Growth with a Tiny MLP

Problem Setup

Goal: Use a small neural network to predict **next-quarter GDP growth** from the **last 4 quarters**.

- Input features at time t :

$$x_t = (\text{growth}_{t-4}, \text{growth}_{t-3}, \text{growth}_{t-2}, \text{growth}_{t-1})$$

- Target:

$$y_t = \text{growth}_t.$$

- Model:

- small MLP with 1 hidden layer,
- trained with MSE loss and Adam optimizer.

- This is **not** a sophisticated macro model — just a **simple demo** of the ML **workflow**.

Building Features from Cleaned Data

Assume we already have df with a growth column (from earlier).

```
1 import numpy as np
2 import torch
3 from torch.utils.data import TensorDataset, DataLoader
4
5 # 1. Get growth as NumPy array
6 g = df['growth'].values.astype(np.float32)
7
8 # 2. Build (X, y) where each X has last 4 quarters
9 X_list = []
10 y_list = []
11
12 for t in range(4, len(g)):
13     X_list.append(g[t-4:t]) # last 4 quarters
14     y_list.append(g[t])      # next quarter
15
16 X = np.stack(X_list)          # shape: [N, 4]
17 y = np.array(y_list)[:, None] # shape: [N, 1]
18
19 # 3. Convert to tensors
```

Train / Test Split and DataLoader

```
1 N = X_tensor.shape[0]
2 train_size = int(0.8 * N)
3
4 X_train = X_tensor[:train_size]
5 y_train = y_tensor[:train_size]
6 X_test = X_tensor[train_size:]
7 y_test = y_tensor[train_size:]
8
9 train_ds = TensorDataset(X_train, y_train)
10 test_ds = TensorDataset(X_test, y_test)
11
12 train_loader = DataLoader(train_ds, batch_size=32,
13                           shuffle=True)
14 test_loader = DataLoader(test_ds, batch_size=32,
15                           shuffle=False)
```

Defining a Tiny MLP in PyTorch

```
1 import torch.nn as nn
2 import torch.optim as optim
3
4 class GrowthMLP(nn.Module):
5     def __init__(self):
6         super().__init__()
7         self.net = nn.Sequential(
8             nn.Linear(4, 16),
9             nn.ReLU(),
10            nn.Linear(16, 1)
11        )
12
13    def forward(self, x):
14        return self.net(x)
15
16 model = GrowthMLP()
17 optimizer = optim.Adam(model.parameters(), lr=0.01)
18 loss_fn = nn.MSELoss()
```

Training Loop

```
1 num_epochs = 50
2
3 for epoch in range(num_epochs):
4     model.train()
5     epoch_loss = 0.0
6
7     for batch_X, batch_y in train_loader:
8         optimizer.zero_grad()
9         preds = model(batch_X)
10        loss = loss_fn(preds, batch_y)
11        loss.backward()
12        optimizer.step()
13
14        epoch_loss += loss.item() * batch_X.size(0)
15
16    epoch_loss /= len(train_loader.dataset)
17    print(f"Epoch {epoch+1:03d} | Train MSE: {epoch_loss:.6f}")
```

Evaluating and Plotting Predictions

```
1 model.eval()
2 with torch.no_grad():
3     y_pred_list = []
4     y_true_list = []
5     for batch_X, batch_y in test_loader:
6         preds = model(batch_X)
7         y_pred_list.append(preds.numpy())
8         y_true_list.append(batch_y.numpy())
9
10    y_pred = np.vstack(y_pred_list).flatten()
11    y_true = np.vstack(y_true_list).flatten()
12
13    import matplotlib.pyplot as plt
14
15    plt.figure(figsize=(10, 4))
16    plt.plot(y_true, label="True growth")
17    plt.plot(y_pred, label="Predicted growth")
18    plt.title("Next-quarter GDP growth: true vs predicted")
19    plt.xlabel("Test observation index")
20    plt.ylabel("Growth")
```

Wrap-Up

What You Should Take Away

- **Tech stack:** OS → Python → Jupyter/VS Code/Cursor/PyCharm → libraries → AI assistants.
- **Python basics:** data types, lists, dictionaries, loops, functions, classes (attributes, methods, self).
- **Data workflow:** acquire, clean, HP-filter, persist, visualize.
- **PyTorch flavor:**
 - tensors, autograd, models as classes, optimizers, DataLoader.
 - tiny MLP predicting GDP growth from past values.
- Next lecture: **AI and ML theory for economists:**
 - loss functions, overfitting, regularization, evaluation, etc.

AI-Assisted Coding: The Research Architect Approach

The AI Era Challenge

What's changing?

AI tools can now write code, debug algorithms, and translate math into implementations.

Question: If AI can code, what's left for economists?

The Shift

Your **attention** shifts from **how to code** to **what to code and why**.

From Coder to Research Architect

Traditional

- Write every line
- Debug syntax errors
- Focus: **how**

Research Architect

- Design algorithms
- Validate economics
- Focus: **what & why**

Three core skills:

1. **Algorithm Design:** Translate theory into computational steps
2. **Strategic Direction:** Guide AI from simple to complex
3. **Critical Validation:** Verify technical + economic correctness

AI Coding Tools: Three Categories

Chatbox	IDE-Integrated	Autonomous
ChatGPT, Claude	Cursor, Copilot	Claude Code, Aider
Algorithm design Debugging Theory questions	Real-time coding Inline suggestions File awareness	Multi-file projects Refactoring Autonomous debug

Recommendation

Primary: Cursor (IDE) + Claude (chatbox)

Advanced: Claude Code (autonomous)

Tool 1: Cursor

VS Code + AI Integration

Key features:

- **Inline completion:** Tab to accept AI suggestions
- **Chat** (Cmd/Ctrl+L): Ask questions with file context
- **Composer** (Cmd/Ctrl+I): Multi-file edits

Use for: Day-to-day coding, quick iterations, debugging

Tool 2: Claude

Conversational AI for Research

Strengths:

- Long context (~200k tokens) — upload papers + code
- Strong economic reasoning
- Interactive artifacts

Use for:

- Algorithm design: “How to discretize state space?”
- Literature: Upload PDF, ask “Implement Section 3”
- Deep debugging with economic interpretation

Best practice: Provide economic context, not just code

Tool 3: Claude Code

Autonomous Terminal-Based Agent

What it does:

- Reads/writes multiple files
- Runs code, observes errors, debugs autonomously
- Works across entire projects

Use for:

- Large refactoring
- MATLAB → Python conversion
- Implementing full algorithms from papers

Key difference: Cursor assists you; Claude Code works autonomously

The Iterative Development Method

Core Idea:

Start simple. Grow complexity systematically.

Mathematical analogy: Homotopy continuation

- Solve easy problem $F_0(x) = 0$
- Define path: $H(x, t) = (1 - t)F_0(x) + tF_1(x)$
- Track solution as $t \rightarrow 1$ to reach hard problem $F_1(x) = 0$

Here: Start with simple model, incrementally add features

Pattern: Algorithm → Extensibility check → Code with diagnostics → Validate → Extend → Repeat

Why This Works

Benefits:

1. Validation at each step

- Know exactly where bugs come from
- Clear success criteria

2. AI succeeds on simple problems

- Build confidence incrementally
- Good starting point for complex versions

3. Extensible by design

- Check algorithm accommodates future features
- Avoid major refactoring later

Key Insight

Design for extensibility from the start, validate continuously

Example: RBC Model Development

Goal: Build stochastic RBC model

Don't start with full model!

Start simple:

- Deterministic RBC
- Log utility, Cobb-Douglas production
- Fixed labor supply
- Has closed-form steady state for validation

Solve via VFI: Iterate on Bellman equation until convergence

Why VFI for deterministic case?

- Algorithm naturally extends to stochastic case
- Same code structure accommodates shocks later

Workflow Step 1: Design Algorithm

Your prompt to AI:

"Design a VFI algorithm for deterministic RBC:

- Production: $Y = AK^\alpha$
- Bellman: $V(k) = \max_{k'} \log(AK^\alpha - k') + \beta V(k')$
- Use grid search over k'

"

AI responds with algorithm

CRITICAL: Before implementing, ask AI:

"Can this algorithm extend to stochastic productivity? Where would shocks enter?"

Refine algorithm: Ensure state space structure accommodates (k, A) later

Workflow Step 2: Implement with Diagnostics

Your prompt to AI:

"Implement this algorithm. Include diagnostic checks:

1. Convergence plot (sup norm vs iteration)
2. Compare $V(k)$, $k'(k)$ to analytical steady state
3. Compute and plot Euler equation residuals

"

AI writes code with built-in validation

Your job:

- Run code
- Check plots: Does $k'(k)$ match steady state formula?
- Euler residuals near machine precision?

Workflow Step 3: Extend Algorithm

After validating deterministic version:

Your prompt to AI:

"Extend algorithm for stochastic productivity:

- AR(1): $\log A_{t+1} = \rho \log A_t + \epsilon_t$
- State space: (k, A) grid
- Bellman: $V(k, A) = \max_{k'} u(c) + \beta E[V(k', A')|A]$
- Discretize using Tauchen method

Show updated algorithm."

Review algorithm before coding

Workflow Step 4: Update Code

Your prompt to AI:

"Implement stochastic version. Add:

1. Simulation of $\{k_t, A_t, Y_t, C_t\}$ for $T = 1000$ periods
2. Impulse response to productivity shock
3. Report: $E[Y]$, $\sigma(Y)$, $\text{corr}(Y_t, Y_{t-1})$
4. Verify: Setting $\rho = 0, \sigma = 0$ recovers deterministic case

"

Validate:

- Check IRF shapes (positive shock → output rises)
- Verify moments make economic sense
- Confirm limiting case works

Workflow Step 5: Iterate

Next iteration: Add variable labor

Repeat the pattern:

1. Extend algorithm: Add labor choice, $MRS = MPL$
2. Check extensibility: Room for labor adjustment costs?
3. Implement with diagnostics: Plot labor supply, verify FOC
4. Validate: Does labor respond to shocks correctly?

Continue: Adjustment costs → Multiple sectors → Financial frictions → ...

The Progression

Each validated iteration becomes foundation for the next

Key Principles of Iteration

1. Algorithm first, code second

- Design and refine algorithm before implementing
- Cheaper to fix algorithm than code

2. Check extensibility immediately

- “Can this accommodate feature X?”
- Avoid painting yourself into a corner

Key Principles of Iteration

3. Build diagnostics into code

- Comparison plots, Euler residuals, limiting cases
- Code should validate itself

4. One feature at a time

- Know which change caused any bugs
- Clear attribution of results

Your Validation Checklist

At every iteration, verify:

Technical

- Convergence?
- NaNs/infinities?
- Grid adequate?

Economic

- Results sensible?
- Comparative statics correct?
- Limiting cases work?

Golden Rule

AI writes code. YOU validate economics.

Debugging with AI

When something goes wrong:

1. Provide full context: code + error + expected behavior
2. AI suggests diagnostics
3. You run diagnostics, report back
4. Iterate until resolved

AI catches:

- Indexing errors, broadcasting issues
- Numerical problems (overflow, underflow)
- Logic errors in algorithms

You catch:

- Nonsensical equilibria
- Wrong comparative statics
- Economic assumptions violated

Best Practices

1. Always start simple

- Find analytically tractable special case

2. Design for extensibility

- Check algorithm accommodates future features

3. Build in validation

- Code should include diagnostic checks

4. Use AI strategically

- AI handles implementation
- You handle algorithm design and validation

5. Validate incrementally

- Never move forward with unresolved issues

The Research Architect Mindset

Traditional	Research Architect
Writes every line	Designs algorithms
Debugs syntax	Validates economics
“How to code X?”	“Does X make sense?”
Bottleneck: coding	Bottleneck: understanding

Your edge:

- Economic theory expertise
- Spotting nonsensical results
- Knowing which simplifications preserve mechanisms

Core Takeaway

AI democratizes **implementation**.

Knowing what to implement and **validating correctness** remain
your unique value.

The workflow:

1. Design algorithm (check extensibility)
2. Implement with diagnostics
3. Validate thoroughly
4. Extend algorithm
5. Repeat

Summary

AI tools: Cursor (IDE) + Claude (chatbox) + Claude Code (autonomous)

Iterative method:

- Start simple (analytical benchmark)
- Design extensible algorithms
- Code with built-in validation
- Add features incrementally

Your role: Algorithm architect + Economics validator

Next: We'll see Python basics and how this workflow looks in practice